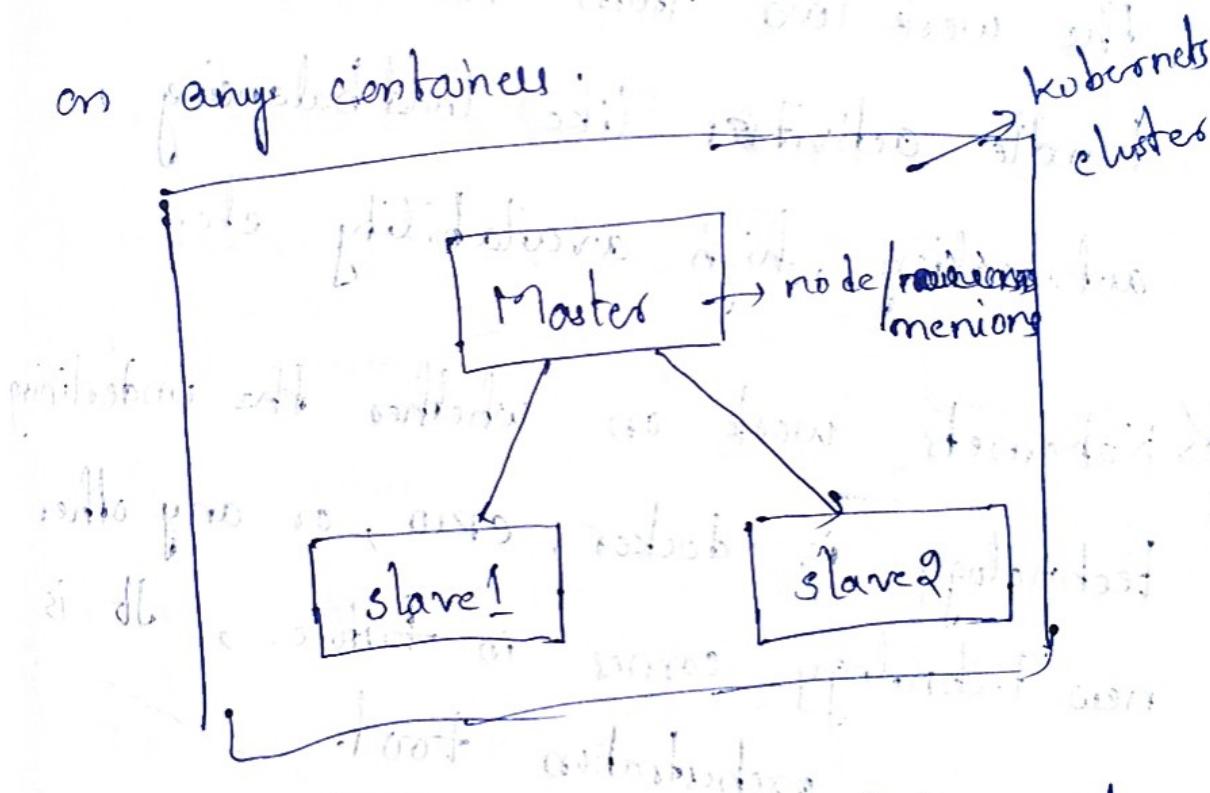


KUBERNETES

→ It is a container orchestration tool, where as docker swarm is docker container orchestration tool.

→ It means docker swarm work on only docker containers, where kubernetes work on any containers.



→ We have one Master machine and slaves in the cluster

→ The individual node is known as
minions

→ The combination of these members is known as kubernetes cluster.

→ Master is the main machine which triggers the containers orchestration. It distributes the work load to the slaves.

→ Slaves are the nodes that accept the work load from the master and handle activities like load balancing, autoscaling, high availability etc.

* Kubernetes work on whether the underlining technology is docker, crio, or any other new technology comes in future, It is a generic orchestration tool.

Kubernetes objects

① Pod :-

- It is the layer of abstraction on top of a container.
- It is the smallest object that kubernetes can work on.
- In the Pod we have containers.
- The advantage of Pod is that kubectl commands will work on the Pod and Pod communicates these instructions to the containers irrespective of which technology containers are in the Pod.
- It means it acts like a translator between containers of any technology with kubernetes commands.



★ → Kubectl commands can work upto the level of Pod, then Pod translates in such a way that underlying container can understand.

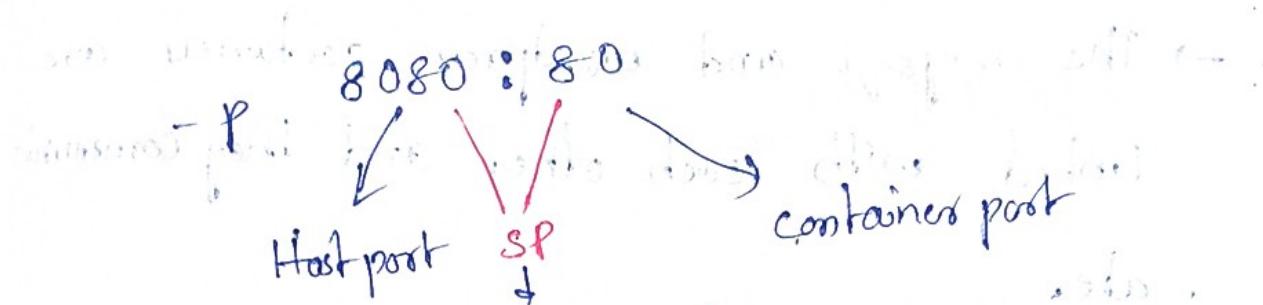
Interview:

Q what is pod? why it is required?

→ Because kubernetes is container orchestration tool its not restricted to docker, it can work on any kind of container it could be docker, criod etc. So they have designed the kubernetes in such a way that kubernetes can communicate upto the level of Pod. It is the responsibility of Pod to understand which kind of container is inside it and according do the necessary translations is called as Pod.

② Service :-

→ It is used for port mapping and n/w load balancing.



→ In between host port and container port

we have service port

→ Service port ensures that same ip address

is assigned to the newly recreated container

(if container crashes, there is a chance

of having another ip address for newly

recreated container, then communication

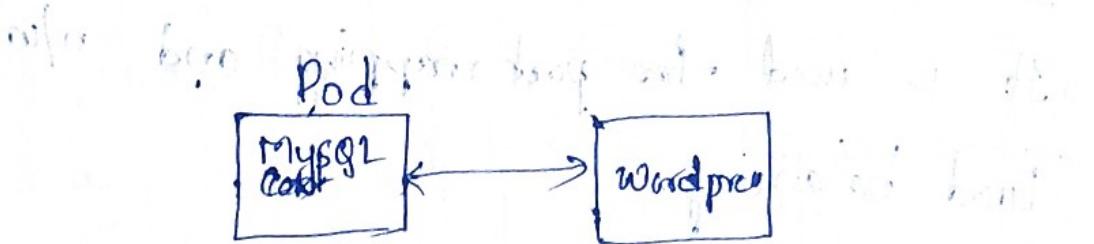
between that container, and other will be

lost). So to avoid that miscommunication

service port taking care of

assigning same ip address to the containers

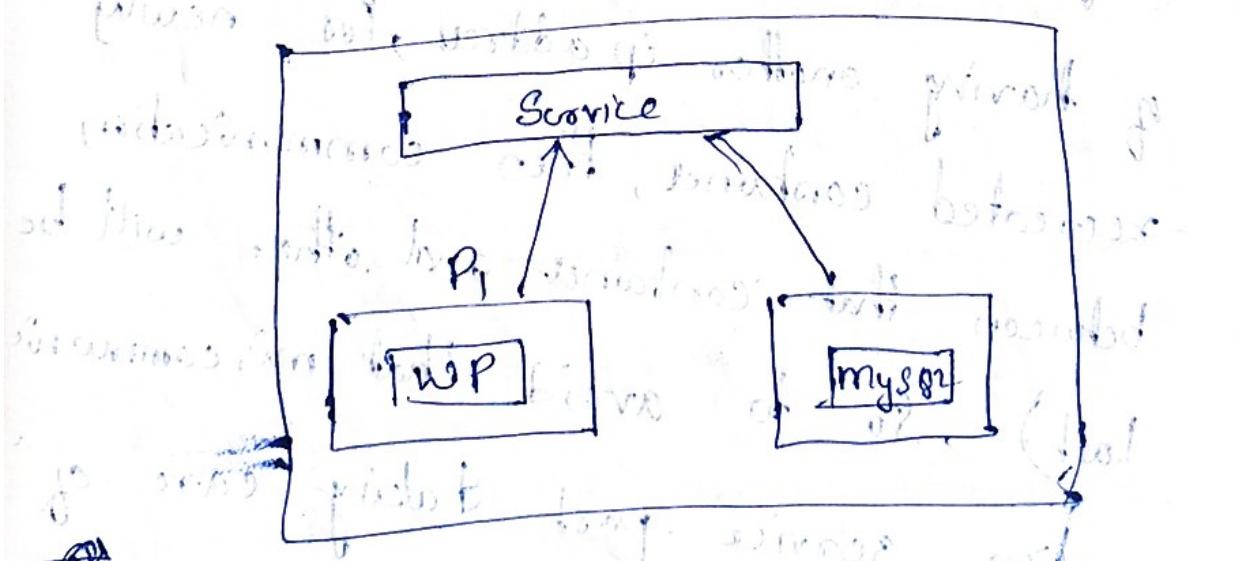
eg:-



- The mysql and wordpress containers are linked with each other and they communicate with each other.

- But in any case if the mysql gets crashed, then kubernetes creates another replica automatically but it comes up with different ip address, hence communication between them is lost.

- To avoid that we have service



- Service is off separately storing the IP and port.

addresses of each Pod, and assigns the same ip address for newly created Pod.
→ whenever same ip address is assigned then communication between ~~contai~~ Pods remain same.

③. Replication Controller :-

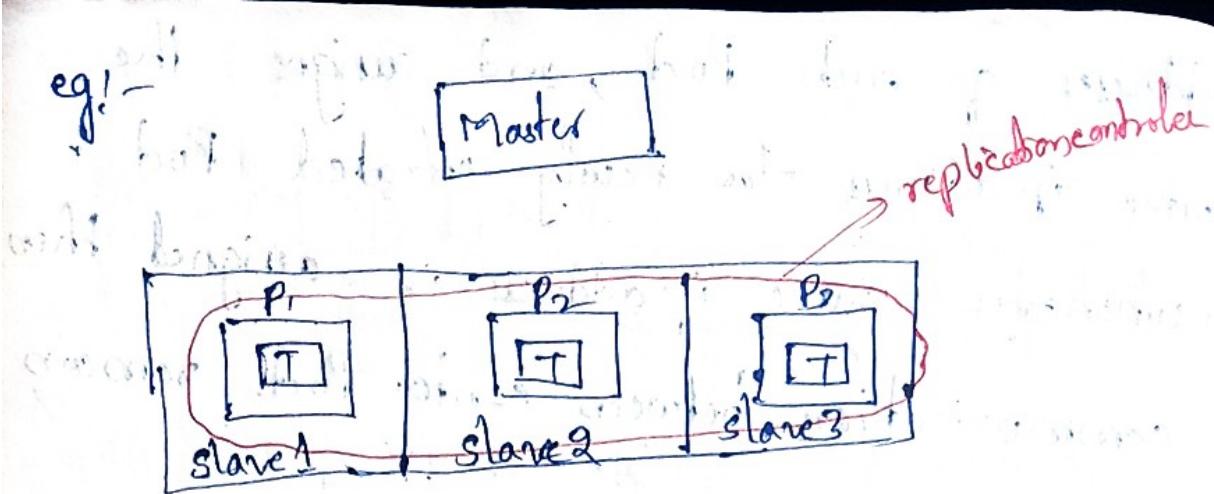
It is used for performing two activities of container orchestrations, they are

- (i) Load Balancing
- (ii) Scaling.

→ This is used for managing multiple replicas of Pods and also performing scaling.

→ To increase the no. of replicas for a particular service we have to increase the no. of replicas by creating more replication controller.

→ To decrease the no. of replicas we have to delete the replication controller.



- We have 3 slaves, on each of them has Pod, in the Pod, tomcat container is there.
- To maintain all these pods on slaves we have another layer called replication controller.
- For load balancing, and scaling of replicas thus use replication controller.

④ Replicaset

- It is also similar to replication Controller but some it has some advanced feature like selector is implemented.

- It is also used for maintaining replicas which are running on different slaves.
- It also performs load balancing and scaling. but it has attribute called selector.
- The selector has ability to pick up particular element based on label and use them.

Eg:-

- We want to create 3 replicas of tomcat.
- But already 2 replicas of tomcat are running.
- Then selector sees that already two tomcat Pods are running and then it creates one Pod of tomcat and reaches the desired state of 3.
- It checks the components which you are creating in the cluster, if they

are available, it will use them otherwise it will create.

→ In this we have a reusability factor.



⑤ Deployment :-

→ It can do load balancing, scaling and also rolling update.

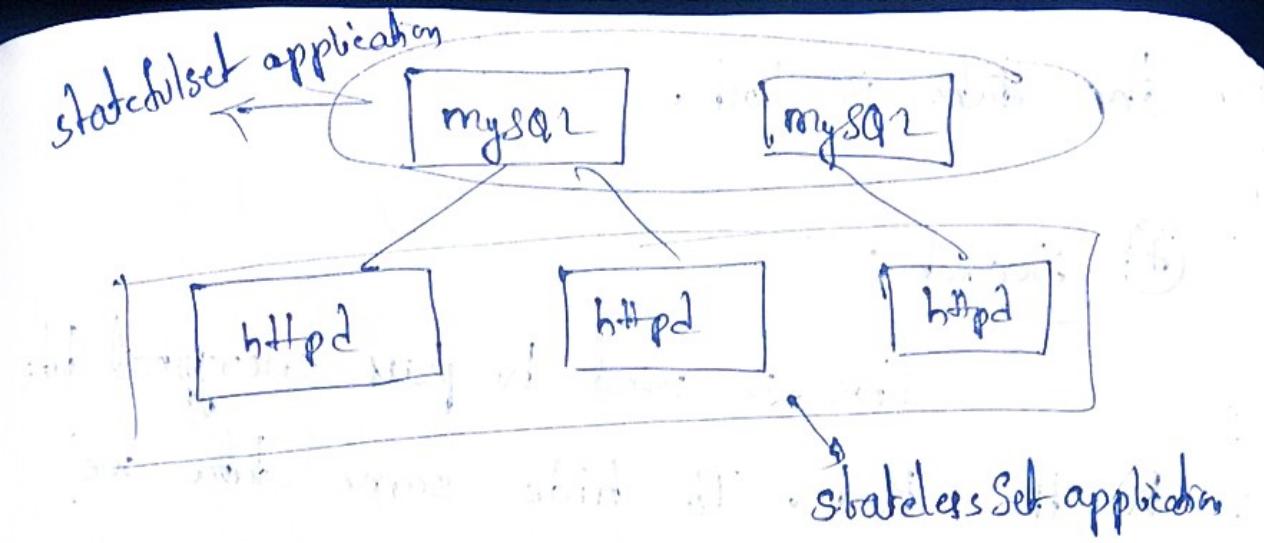
⑥ StatefulSet :-

These are used to handle stateful applications like database where that data has to be preserved even if the container got crashed.

→ The newly created container must come up with the previous data.

→ For this purpose we use statefulSet.

e.g:- All databases



→ The httpd act as user interface; it takes the data and pass it to mysql database. So these are called statelessSet applications, where mysql databases are called statefulSet applications.

Interviewer: How you maintain databases in docker

→ We don't have databases in the docker swarm cluster, but we are maintaining database with regular installations and linking with docker swarm cluster because if the database is maintained in the swarm cluster, then when the container crashes

the data is lost.

⑦ Secret :-

This is used to pass encrypted data to the Pod. To hide some data we use secrets.

⑧ Ingress :-

The mapping of the domain name with sub domains is done with ingress.

eg:- amazon.in / something → sub domain
↳ main domain

⑨ Horizontal Pod Autoscaler :-

→ The no. of pods will automatically increase or decrease depending on the load of the application.

→ In docker swarm we don't have autoscaling because it's built with pods.

⑩ Persistent Volume :-

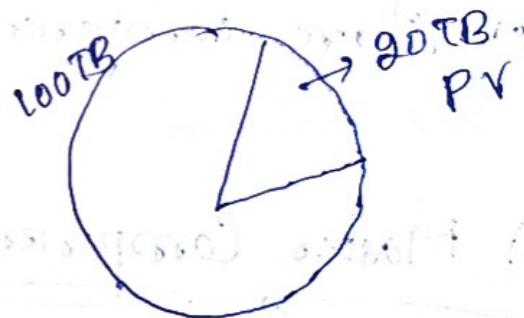
It is used to specify the section of storage that should be used for volumes.

⑪ Persistent Volume Claim :-

It is used to reserve a certain amount of storage for pod from the persistent volume.

→ These two are used for backup mechanism.

e.g:-



Let's say from 100TB of cluster, 20TB is reserved for volumes, it is called PV.

Now I want to use 1TB for creating a container from 20TB of PV, then it

is called PVC.

⑫ Namespace! -

It is used for creating partitions in the cluster. Pods are running in a namespace. Pods cannot communicate with other Pods running in other namespace.

Master Slave Components / Master Slave architecture

① Master components

② Slave components

① Master Components :- It is responsible for managing the execution and lifecycle of containers within the K8s.

a) Container runtime :- It is generally a docker running on the master machine. We call it container runtime.

call that in kubernetes as container

runtime.

→ don't use docker word here.

⑥ Kube apiserver :-

- It acts as a validator.
- It is a gateway to the cluster.
- It works as a gatekeeper for authentication and it validates if a specific user is having permissions to execute a specific command.

e.g. → If he has permissions, then it passes the ~~request~~ to scheduler whether he has permission to delete a cluster, then I am trying to delete a cluster, then I executed delete command for that cluster. Then Kube apiservice checks whether he has permission to delete or not. If I have permission its passes.

To the scheduler who performs the operation which I want to do.

↳ ~~which operation is performed by which component with permission~~

③ Kube scheduler:-

- This process accepts the instruction from apiserver after validation and starts an application on a specific node or set of nodes.
- The hardware resources are estimated by kube scheduler.
- e.g:- I want to create 3 tomcat replicas then kube scheduler first checks how many slaves we have, how much amount of storage is present on each slave, and on which machine to create these replicas, all these are estimated by kube scheduler.

④ Control manager or Controller:-

- Maintaining the ~~desired~~ desired

count of replicas is done by this controller.

e.g. let's say I have desired count of

6 replicas, suddenly one replica crashes. Then control manager automatically creates replicas and maintains desired count of 6.

etcd :-

- It is a repository or database.
- It stores all the information about the cluster i.e. how many slaves are there, no. of replicas, ip addresses, every information about the cluster.
- It is the core repository, the entire information of a cluster is stored in etcd, it makes a db layer.

It stores key-value pairs in JSON format.

Slave Components:-

a) Containers runtime:-

Some containerization technology should be there; generally docker is used.

b) Kubelet:-

This is the actual process that takes the orders from the scheduler and deploy an application or any operation on a slave.

→ It is present on both master and slave.

c) Kube proxy:-

→ It will take the request from services to pod.

→ It has the intelligence to forward a request to a near by pod.

e.g. - If an application pod wants to

communicate with db pod then kube proxy will take that request to the near by pod. So that latency will be reduced.

Kubernetes installations → These are divided into two types, they

Unmanaged / self managed → Kubernetes setup

1) KOPS

2) kubeadm

3) KIND

Managed → Kubernetes setup

1) EKS (AWS)

2) GKE (GCP)

3) AKS (Azure)

→ In managed kubernetes setup, the cloud automatically setup the cluster where as in self/unmanaged we manually provision it.

have to setup their cluster in AWS
by getting all the things ready.

① KOPS

~~CloudFormation~~ is used for creating the cluster.

→ It uses S3, in AWS, it is like a drive for storage, they call it as bucket.

→ Route53 → It acts like DNS (it converts the domain name into ip address)

→ Auto scaling → It also activates auto scaling, i.e if a server gets crashed, it automatically creates a server in the cluster.

→ If the KOPS server wants to activate all these services, it needs some permissions given by IAM.

S3
EC2
VPC
Route53
Autoscaling

Step - I :-

- Launch Linux EC2 instance in AWS
- Create and attach IAM role to EC2 instance

Step - II :-

- Install kops on EC2.

```
curl - LO https://github.com/kubernetes/kops/releases/download/`curl -s https://api.github.com/repos/kubernetes/kops/releases/latest | grep tag-name | cut -d '"' -f 4`/kops-linux-amd64
```

- Now change execute permissions of kops -

linux-amd64

```
chmod +x kops-linux-amd64
```

- Now move that file into /usr/local/bin/kops

```
sudo mv kops-linux-amd64 /usr/local/bin/kops
```

Step - III :- Install kubectl

```
curl -LO https://storage.googleapis.com/kubernetes-release/release/stable.txt | curl -s
```

```
https://storage.googleapis.com/kubernetes-release/release/stable.txt) /bin/linux/amd64/kubectl
```

→ change permissions of kubectl

```
chmod +x ./kubectl
```

→ Now move that file in /usr/local/bin

```
sudo mv ./kubectl /usr/local/bin/kubectl
```

Step - IV :- Create S3 bucket in AWS

```
aws s3 mb s3://name of bucket in k8s
```

→ region us-east-1

→ Give any name for the bucket, it must be unique

→ give the region, in which regions your kops servers present.

step 5:- Create private hosted zone in AWS Route 53

Head over to our Route 53 and create AWS Route 53 hosted zone

choose name for example (sudarshans.in)

→ Now choose type as private hosted zone

for VPC

→ Select default VPC in the region you are

getting up your cluster.

→ Hit Create.

step 6:- Now configure environment variable

open .bashrc file

vim ~/.bashrc

Add following contents into .bashrc,

you can choose any arbitrary name for cluster and make sure bucket name matches the one you created in previous step.

`export KOPS_CLUSTER_NAME=project-in`

`export KOPS_STATE_STORE=s3://your-bucket-name-in-k8s`

→ Then run command to reflect variable added to .bashrc

`source .bashrc`

step-VII!: Create ssh key pair

→ This key pair is used for ssh into kubernetes cluster

`ssh-keygen`

step-VIII!:

Create a kubernetes cluster

```
kops create cluster \n\n--state=${{KOPS_STATE_STORE}} \n\n--node-count=2 \n\n--master-size=t3.medium \n\n--node-size=b3.medium \n\n--zones=us-east-1a \n\n--name=${{KOPS_CLUSTER_NAME}} \n\n--dns private \n\n--master-count=1
```

step-1!: Create kubernetes cluster

```
Kops update cluster --yes --admin
```

step-2!: Execute the command now and
then until the cluster gets created.

```
Kops validate cluster
```

step-3!: To delete the cluster permanently

```
Kops delete cluster --yes
```

Step - III To connect to the master

ssh admin@api.javahome.in

We can change number of nodes and no. of masters using following command.

kops edit ig nodes change minsize and maxsize to 0

kops get ig - to get master node name

kops edit ig - change min & max size to 0

kops update eluster -zyes

③ KIND:-

step - I :-
→ Install docker

step - II :-

Install kubectl

→ go to kubectl install site and select
for which flavor you are using. Then
copy and paste those commands

step - III :-

Install kind software

→ go to kind install site
copy those commands and paste

step - IV :-

It requires some configuration file

copy those commands in vim

`vim config.yaml`

step - V :- Now create cluster

`kind create cluster --name mycluster --config config.yaml`

Managed Kubernetes setup

① EKS (Elastic Kubernetes Service)

Step - I :-

First launch Linux service and give IAM permission bind attach, if to EC2 members.

Step - II:-

→ first install eksctl

→ go into google and select eksctl inball then inball it.

Step - III:-

→ Now inball kubelet

Step - IV:- → To activate the cluster or to create the cluster

eksctl create cluster

--region us-east-1
--node-type t3.medium

--nodes 3

--name new-cluster

step - V :-

eksctl create cluster --name my-cluster --region us-east-1

eksctl delete cluster --name my-cluster --region us-east-1

step - VI :-

To see the cluster search on AWS or
AWS Lambda's marketplace of choices.

step - VII :-

Now install certain plugins from the

site

<https://docs.aws.amazon.com/eks/latest/usguide/install-aws-iam-authenticator.html>

(or)

when type kubernetes command, it will show
the website url.

② GKE (Google Cloud Kubernetes Engine)

Step - I: Create google cloud account.

and then go to console.

Step - II:

Then go into navigation symbol
click on kubernetes engine → cluster →
switch to standard cluster → create

Step - III:

If you get error as insufficient regional

quota, then go to default-pools

nodes → in that give 75, where

the no. is showing 100, then enter

create cluster

① To see the list of all the nodes

`kubectl get nodes`

② To see info about the nodes in wider o/p

`kubectl get nodes -o wide`

③ To see the detailed info about a node

`kubectl describe nodes node-name`

④ To create a pod

`kubectl run --image image-name v pod-name`

⑤ To see the list of all the pods that have created

`kubectl get pods -o wide`

⑥ To see the info of pods in wider o/p

`kubectl get pods -o wide`

⑦ To get detailed info about a pod

`kubectl describe pods pod-name`

⑧ To go into the bash shell of a pod

kubectl exec -it pod-name -- bash

⑨ To delete a running pod

kubectl delete pod pod-name

There is no delete option in kubectl

* Create a mysql pod

kubectl run --image mysql-mydb:

env MYSQL_ROOT_PASSWORD=sudarshan

→ We have yaml file in the kubernetes

These files are called as definition files

For obvious reasons, we have top level or manifest files

elements in this & every file of a

manifest file has kind: Pod

kind: Pod

metadata

kind: → it represents which type of object you are using eg:- pod, service, deployment, etc.

apiVersion: — Here we have some api

versions for different objects.

These are the resources of K8s cluster

kind	apiVersion	Interview
① Pod	v1	
② Service	v1	
③ Namespace	v1	
④ Secrets	v1	
⑤ Persistentvolume	v1	
⑥ Persistentvolume claim	v1	
⑦ Horizontal Pod AutoScaler	v1	
⑧ Replicaset	apps/v1	PDS = appset
⑨ Deployment	apps/v1	
⑩ Statefulset	apps/v1	

Metadata: if it is data about the data.

Spec :- The containers information is given in this module.

e.g. - If I am creating a Pod, then in the spec, section will give container information.

Labels :-

You can give any data about the object you are creating. In this key and value are not predefined.

e.g. - author : intelligentworkersof

key

value

You can use any keys and any values in the label section.

language

framework

(*)

To run the definition file

kubectl apply -f file-name

* Create a nginx pod

--
apiVersion: v1

kind: Pod

metadata:

name: nginx-pod

labels:

author: intelligible

type: proxy

Spec:

containers:

- name: mynginx

image: nginx

...
highlights in red

kubectl apply -f pod-definition.yaml

→ To delete the pod which is created using definition file

kubectl delete -f definitionfile-name

* Create the postgres database

apiVersion: v1

kind: Pod

metadata:

name: postgres-pod

Labels:

type: db

author: sudarshan

Spec:

containers:

- name: mypostgres

image: postgres

env:

-- name: POSTGRES_PASSWORD

value: intelligent

-- name: POSTGRES_USER

value: myuser

- name: POSTGRES_DB
value: mydb

④ Create a jenkins pod

--
apiVersion: v1

kind: Pod

metadata:

name: jenkins-pod

labels:

author: Sudarshan

type: CI-CD implementation

Spec:

containers:

- name: myjenkins

image: jenkins/jenkins

ports:

- containerPort: 8080

hostPort: 8686

→ To create .yaml files of kubernetes in pycharm.

Step ①:-

Download a "kubernetes and openshift resource support" plugin

Step ②:-

Now open pycharm, go to file
file → settings → plugins → click on gear
icon → click on → install plugin from disk
then go to download section and select the
downloaded plugin

Namespace :-

We have different types of namespaces in the kubernetes setup. To see the namespaces

[kubectl get namespace]

- In this we a namespace called default
- The pods which are created will be running on default namespace. The namespace is nothing but making the partition in the cluster.

- We can create our own namespaces.

vim definitionfile.yaml

--
apiVersion: v1

kind: Pod Namespace

metadata:

name: test-ns

To run our pod in the "test-ns".

vim definitionfile.yaml

--
apiVersion: v1

kind: Pod

metadata: ghost-pod

vim definitionfile.yaml

- apiVersion: v1

kind: Pod

metadata:

name: ghost-pod

namespace: test-ns

Labels:

type: cms

author: intelligent

Spec:

containers:

- name: ghost

image: ghost

env:

- name: NODE_ENV

value: development

...

To check the customised namespace

kubectl get pods -n namespace-name

eg! + kubectl get pods -n test -n

- Q) The command which shows all the pods in all the namespaces

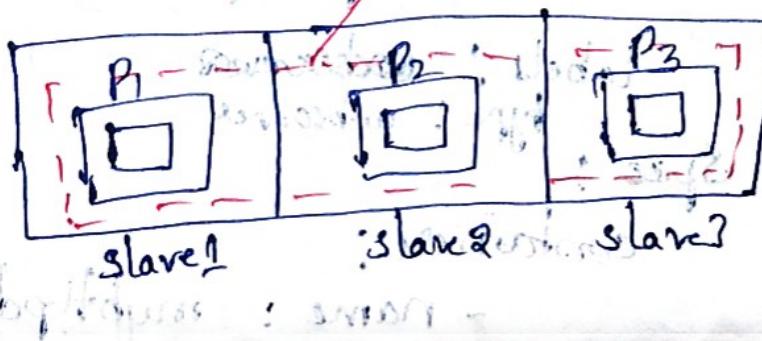
kubectl get pods --all-namespaces

Replication Controller →

- It can perform two activities, they are (i) load balancing (ii) scaling.
- The layer which controls/maintains all the replicas in a cluster is called Replication Controller
- In the Replication Controller we have Pods

In the pods we have containers

↳ Replication controller



template:- If it is just like blueprint, i.e all the replicas created, structure must be same i.e., maintaining the structure of all the replicas same.

If has two child elements, they are

metadata and spec

yaml definition file.yaml

```
apiVersion: v1
kind: ReplicationController
metadata:
  name: httpd-replica
  labels:
    type: webservice
    author: infellight
```

Spec:

```
replicas: 3
```

template:

```
  metadata:
    name: httpd-pod
```

```
    labels:
```

```
      type: webservice
```

```
    Spec:
```

```
      containers:
```

```
        - name: myhttpd
```

image: httpd
ports:
- containerPort: 80
hostPort: 8080

kubectl get all

it shows all the pods information.

Replicaset:— It performs the same actions of replicationController, but it has another element i.e. selector, the selector first it checks whether the replicas are present or not,
→ If the replicas are present, it uses them and creates and achieves the desired count.

→ It is mainly used for Reusability.

Syntax
selector:
matchLabels:
type:

vim definitions.yaml

apiVersion: apps/v1

kind: ReplicaSet

metadata:

name: tomcat-rls

labels: app: tomcat-rls
type: appserver

author: intelligent

Spec:

replicas: 2

selector:

matchLabels:

type: appserver

template:

metadata:

labels: name: tomcat-pod

labels:

type: appserver

Spec:

contains:

- name: mytomcat

image: tomcat:8080

ports:
- containerPort: 8080

- hostPort: 8181

Scaling can be done in two ways

- ① open your definition file, in that change the replicas number, then run the command as

kubectl replace -f definitionfile.yaml

- ② Another way is from the command line itself we can scale

kubectl scale --replicas=2 -f definitionfile.yaml

g! kubectl scale --replicas=2 -f definitionfile
2.yaml

To delete all the pods in the cluster

kubectl delete all

g! kubectl delete all

g! kubectl delete all

g! kubectl delete all

Deployment :-
It is absolutely same as replicaset, but it can perform rolling update
→ It can perform Replicaset activities too

Yml definition file .yml

apiVersion : apps/v1

kind : Deployment

metadata :
name : nginx-deployment

Labels : type=nginx

type : proxy

author : intelligent

Spec :

replicas : 3

selector :

matchLabels : type=nginx

type : proxy

template :

metadata : In, object

name : nginx-pod

Labels :

type : proxy

author: intelligent

Spec:

containers:

- name: mynginx

image: nginx

ports:

- containerPort: 80

hostPort: 9090

DeploymentSet:

It is used to set no. of replicas = no. of slaves, i.e. how many slaves are

there, that many replicas must be created

automatically.

→ Deployment definition file and deployment definition

file are same but in deployment file we

won't give replicas count.

→ It automatically checks how many slaves.

→ It automatically checks if there are enough slaves that becomes the desired

count of slaves.

vim definitionfile.yaml unter k8s

apiVersion: apps/v1
kind: Deployment
metadata:
 name: httpd-deamonset
 labels:
 type: webserver

Spec:

selector:
 matchLabels:
 type: webserver
 template:
 metadata:
 name: httpd-pod
 spec:
 containers:
 type: webserver

Spec:

containers:
 name: myhttpd
 image: httpd
 ports:
 - containerPort: 80
 hostPort: 8888

..

Scenario :- The deamonset is used, when you want to have log collectors on each server. i.e., the copy of the log collecting software/applications must present on each server, in this scenario we use deamonset.

Secret :-

This is used to encrypt the password or any data.

Syntax :-

```
type: Opaque  
stringData:  
a: your password/any data
```

→ Now ~~create~~ create a secret definition file and use that data or variable in other definition files in which you want to encrypt.

* Create a mysql database and the password must be encrypted or not directly mentioned in the definition file.

→ First create Secret definition file vim definitionfile.yaml

apiVersion: v1

kind: Secret

metadata:

name: mysql-secret

labels

author: intelligit

type: Opaque

stringData:

a: intelligit

→ Now pass this variable 'a' in the definition file of database

```
vim definitionsfile.yaml  
---  
apiVersion: v1  
kind: Pod  
metadata:  
  name: mysql-pod  
  labels:  
    type: db  
    author: intelligit  
  
spec:  
  containers:  
    - name: mydb  
      image: mysql  
  env:  
    - name: MYSQL_ROOT_PASSWORD  
      valueFrom:  
        secretKeyRef:  
          name: mysql-secret  
          key: a
```

→ We can also encrypt the password intelligit

"echo "intelligit" | base64"

→ And change the permission of secret definition file, then group and others won't even touch the file.

④ Create a secret file for postgres database
vim definitionsfile.yaml

apiVersion: v1

kind: Secret

metadata:

name: postgres-secret

labels:

author: intelligent

type: Opaque

stringData:

password: intelligent

username: myuser

dbname: mydb

→ Instead of using a, b, c, use some meaningful names as above

Now create postgres deployment file.

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: postgres-deployment
  labels:
    type: db
    author: intelligent
spec:
  replicas: 2
  selector:
    matchLabels:
      type: db
  template:
    metadata:
      name: postgres-pod
      labels:
        type: db
    spec:
      containers:
        - name: postgres-mydb
          image: postgres:11-alpine
          env:
            - name: POSTGRES_PASSWORD
```

- valueFrom : ~~password~~ ;
secretKeyRef :
name : postgres - secret
key : password
- name : POSTGRES - USR
valueFrom :
. secretKeyRef :
name : postgres - secret
key : username
- name : POSTGRES - DB
valueFrom :
secretKeyRef :
name : postgres - secret
key : dbname

...
→ To decode the password which is encrypted, search in google & co "base64 to string", then paste the code, you will get actual password.

* Service :-

It is used for portmapping and making your application externally available (ie we can access from browser)

→ It is classified into 3 types

① NodePort

② Load Balancer

③ ClusterIP

(i) Headless

* ① NodePort :-

It is used for network load balanc-

-ing.

example:-

let say your ~~host~~ pod is running

on the 1st slave, you can access that

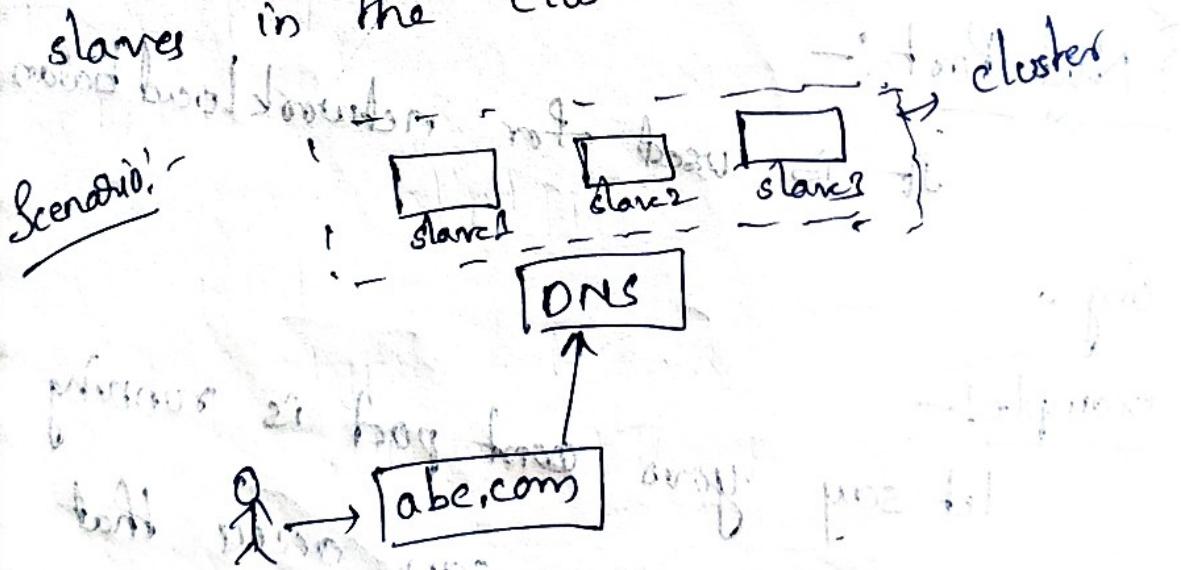
application using 1st slave ip address,

you can't access with the ip address
of other slaves which are in the cluster.

→ So when you use nodeport, you can access from other slaves also, this is called as Network Load balancing.

② Load Balancer :- (also known as load balancing)

It will generate one ip for the entire cluster, no one registers that one ip in the DNS, it will work when you delete some slaves or add some slaves in the cluster.



For the entire cluster generate one address and register in the DNS server.

→ This kind of feature is only available in

kubernetes only, and that too on managed
kubernetes

③ Clusterip :-

It is used when some pods in the cluster which don't want to communicate with external world but they have to communicate with other pods in the cluster internally.

Scenario!

lets say we have Wordpress and MySQL
Now wordpress can communicate with MySQL
and vice versa. but customer can access
only wordpress but not the MySQL database
→ The database access can't be given to
external world , in this scenario we
use clusterip.



① NodePort :-

- * Create a service definition file for port mapping of httpd Pod.

→ first create httpd Pod definition file and then create service definition file based on the labels given in the definition file.

vim definitionfile.yaml

```
apiVersion: v1
kind: Pod
metadata:
  name: httpd-pod
  labels:
    type: webserver
    author: intelligent
spec:
  containers:
    - name: myhttpd
      image: httpd
```

...

vim@service1:~\$ kubectl get svc

apiVersion: v1

kind: Service

metadata:

name: httpd-service

labels:

author: intelligent

Spec:

type: NodePort

ports:

- targetPort: 80

port: 80

nodePort: 30001

selector:

author: intelligent

type: websocket

→ targetPort:— It is nothing but containerPort

→ nodePort:— It is nothing but hostPort

→ port:— It is service port

→ The service files make changes based on the labels given in the selector section.

→ To get the list of services.

`kubectl get svc`

Service port: ←

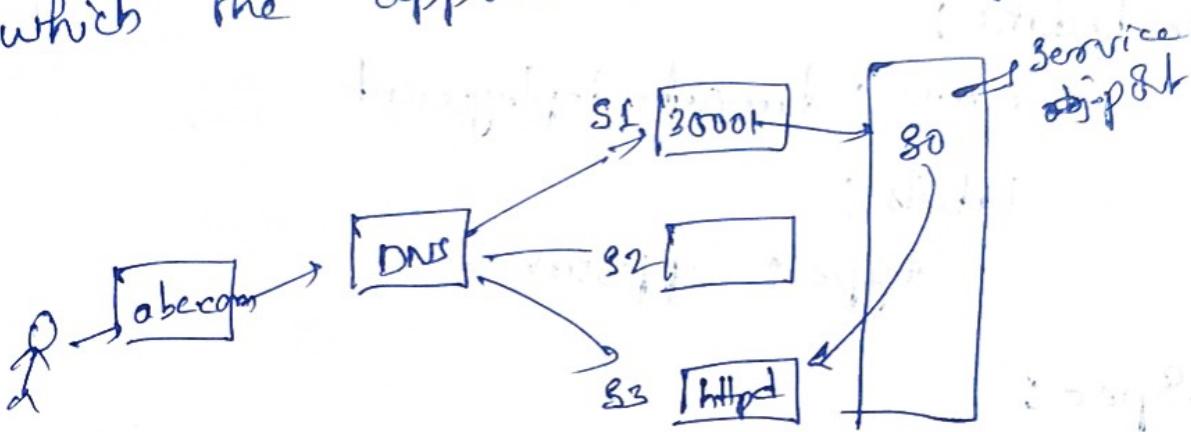
It is the common port mapping between the DNS and the slaves.

Let's say, we have 3 slaves linked with DNS, and my httpd application is running on 2nd slave, if the DNS is sending request to 3rd slave, then httpd will respond, or if it is sending request to 1st and 2nd slave, then it creates a problem.

→ To overcome that problem we are assigning

it with service port.

→ Then service port takes the request from DNS, and sends it to the slave on which the application is running.



→ nodePort range starts from 30,000 to 32767

→ If you give nodePort or it will assign that number

→ If you don't give nodePort number, it will automatically do port mapping

nodePort (hostPort)

* Create a tactical Deployment and service file for it.

vim definitionfile.yaml

```
apiVersion: apps/v1
kind: Deployment
metadata:
```

```
  name: tomcat-deployment
```

```
  labels:
```

```
    type: appserver
```

```
Spec:
```

```
  replicas: 2
```

```
  selector:
```

```
    matchLabels:
```

```
      type: appserver
```

```
  template:
```

```
    metadata:
```

```
      name: tomcat-pod
```

```
    labels:
```

```
      type: appserver
```

```
Spec:
```

```
  containers:
```

```
    name: tomee
```

```
    image: tomee
```

vim .service.yaml

apiVersion: v1

kind: Service

metadata:

name: tomcat-service

labels:

type: appservice

Spec:

type: NodePort

ports:

- targetPort: 8080

port: 80

nodePort: 30002

selector:

type: appservice

→ Deploy both the files in the cluster

and access from any ip address of

any slave, application will respond. This
is called network load balancing

② Load Balancer! - It assigns one unique ip for entire cluster; with that ip we can access the application from browser.

* Create a nginx Pod and service file

vim definitionfile.yaml

apiVersion: v1

kind: Pod

metadata:

name: nginx-pod

labels:

author: intelligent

type: proxy

Spec:

containers:

- name: mynginx

image: nginx

multiple containers: This container will run alongside another container.

vim service3.yaml

```
--  
apiVersion: v1  
kind: Service  
metadata:  
  name: nginx-service  
  labels:  
    type: proxy  
    author: intelligent  
  
spec:  
  type: LoadBalancer  
  ports:  
    - targetPort: 80  
      port: 80  
      nodePort: 30008  
  selector:  
    type: proxy  
    author: intelligent
```

③ Clusterip :-

* Create postgres Pod, see that it should not be accessed by the external world.

→ clusterip is the default service of kubernetes cluster.

vim definitionfile.yaml

apiVersion: v1

kind: Pod

metadata:

name: postgres-pod

labels:

type: db

author: intelligent

Spec:

containers:

- name: mydb

image: postgres

...

vim service4.yaml

--
apiVersion: v1

kind: Service

metadata:

name: postgres-service

labels:

type: db

author: intelligit

Spec:

ports:

- targetPort: 5432

port: 5432

selector:

type: db

author: intelligit

→ We are not giving access to external world, so we are not mentioning nodePort,

Kompose :- (docker compose + kubernetes = Kompose)

→ Create one docker compose file and apply the command as

kompose convert

It will convert that file into all necessary definition files.

Install Kompose :-

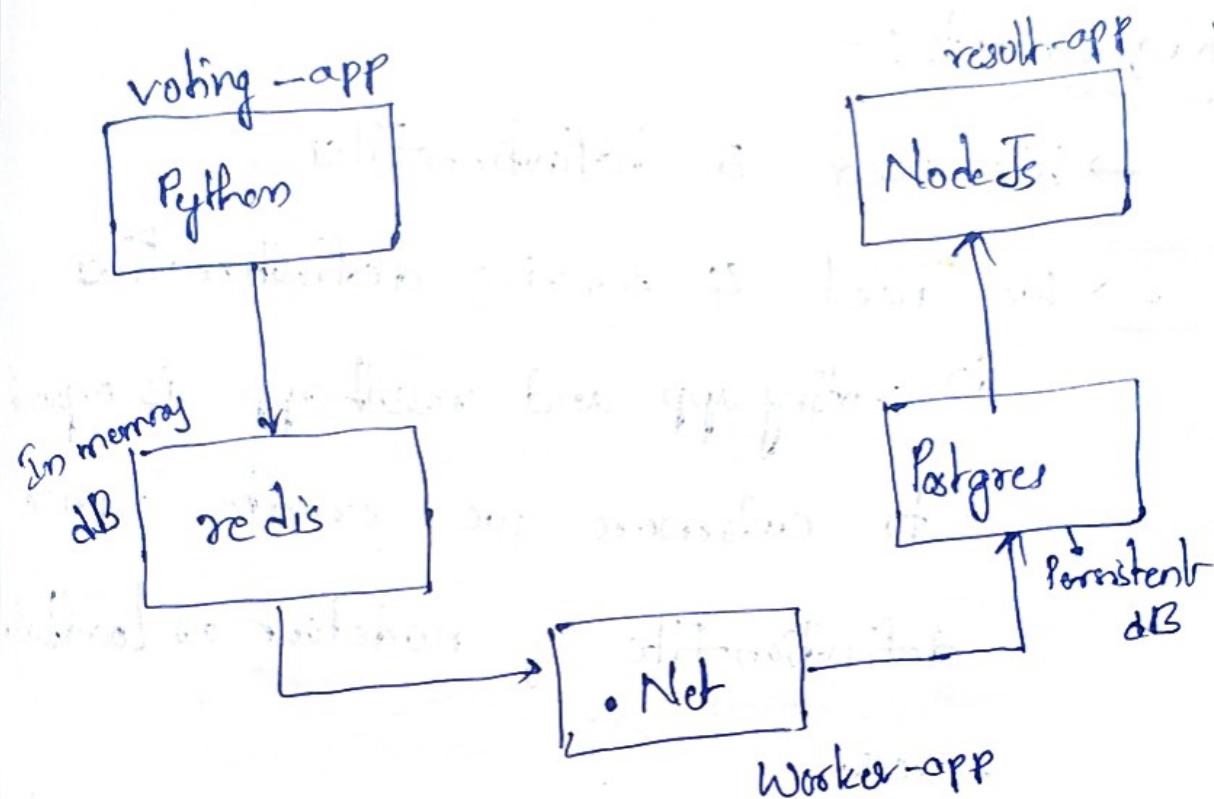
Search in google and install kompose

e.g:- Design docker compose file for mysql and word press and convert into kubernetes definition files

kompose convert

Project :- Voting system

Architecture :-



→ voting app is created by pythonjango
→ voting app is used to cast their
vote work.

→ redis , is used to collect the data
and it pass through the worker app

→ Worker app will filter the data
and send it into postgre db

→ Then postgre stores the data permanently

→ Result-app is created by NodeJs,
and is used to view the results.

Requirements :-

→ We need 5 definition files

→ We need 4 service definition files

(i) voting app and result app is exposed

to customers we create service

definition file of nodePort or Loadbal-

-ancer

(ii) redis and postgres are not exposed

to customers and they are used to

communicate internally, so create

2 clusterip service files.

→ For worker, it is just filtering the
data and passing through to postgres

it is not exposed to customers.

→ create a folder

`mkdir project`

→ Inside the folder : create 5 deployment

definition files and 4 service definition files

→ vim voting-app-deployment.yaml

apiVersion: apps/v1

kind: Deployment

metadata:

name: voting-app-deployment

labels:

name: voting-app

author: intelligent

Spec:

replicas: 2

selector:

matchLabels:

name: voting-app

template:

metadata:

name: voting-app-pod

labels:

name: voting-app

Spec:

containers:

- name: my-voting-app

- image: dockersamples/example-votingapp-vote

...
old-vote-app-service

→ vim result-app-deployment.yaml

apiVersion: apps/v1

kind: Deployment

metadata:

name: result-app-deployment

labels:

name: result-app

author: intelligent

Spec:

replicas: 2

selector:

matchLabels:

name: result-app

template:

metadata:

- name: result-app-pod

Labels:

name: result-app

Spec:

containers:

app - name: my-result-app

image: dockersamples/example

votingapp-result

③ → vim redis-app-deployment.yaml

apiVersion: apps/v1

kind: Deployment

metadata:

name: redis-app-deployment

Labels:

name: redis-app

author: intelligent

Spec:

replicas: 1

selector:

matchLabels:

name: redis-app

template:

metadata:

name: redis-app-pod

Labels:

name: redis-app

Spec:

containers:

- name: my-redis-app

image: redis

④ → vim postgres-app-deployment.yaml

apiVersion: apps/v1

kind: Deployment

metadata:

name: postgres-app-deployment

Labels:

name: postgres-app

author: intelligent

Spec:

replicas: 1

selector:

matchLabels:

name: postgres-app

template:

metadata:

name: postgres-app-pod

Labels:

name: postgres-app

Spec:

containers:

```
- name: my-postgres-app
  image: postgres:13.3
  env:
    - name: POSTGRES_PASSWORD
      value: intelligit
    - name: POSTGRES_USER
      value: myuser
    - name: POSTGRES_DB
      value: mydb
```

⑤ vim worker-app-deployment.yaml

```
---
apiVersion: apps/v1
kind: Deployment
metadata:
  name: worker-app-deployment
  labels:
    name: worker-app
    author: intelligit
spec:
  replicas: 1
```

```
selector:
  matchLabels:
    name: worker-app

template:
  metadata:
    name: worker-app-pod

  labels:
    name: worker-app

spec:
  containers:
    - name: my-worker-app
      image: dockersamples/
example/votingapp-worker
```

...
→ Now create service files for voting app and resultapp.

① vim voting-app-service.yaml

```
apiVersion: v1
kind: Service
metadata:
  name: voting-app-service
```

labels :

~~author~~ :

author : intelligit

Spec :

type : LoadBalancer

ports :

- targetPort : 80

port : 80

nodePort : 30008

selector :

name : voting-app

② vim result-app-service.yaml

apiVersion : v1

kind : Service

metadata :

name : result-app-service

Labels :

author : intelligit

Spec :

type : LoadBalancer

ports :

- targetPort : 80

port : 80

nodePort: 30009

selector:

name: redis-app

③ vim redis-app-service.yaml

apiVersion: v1

kind: Service

metadata:

name: redis-app-service

labels: app: redis-app

author: intelligent

Spec:

ports:

- targetPort: 6379

port: 6379

selector:

name: redis-app

④ vim postgres-app-service.yaml

apiVersion: v1

kind: Service

metadata:

name: postgres-app-service

Labels:

version:

author: intelligit

Spec:

ports:

- targetPort: 5432

port: 5432

selector:

name: postgres-app

...

→ Now deploy all the 9 files and
access from the browser, via

clusterIP: nodePort

~~graph LR~~ > ~~graph LR~~

→ For the above code, for all the files

just create a dockercompose file and

convert into kompose, it will create all
9 definition files

vim myapp.yml

version: '3.8'

services:

my-voting-app:

image: dockersamples/example

votingapp-vote

ports:

- 5050:80

deploy:

replicas: 2

redis-app:

image: redis

ports:

- 6379:6379

worker-app:

image: dockersamples/examplevotingapp-worker

postgres-app:

image: postgres

ports:

- 5432:5432

environment:

POSTGRES_PASSWORD: sudarshan

POSTGRES_USER: myuser

POSTGRES-DB : mydb

result-app :

image : dockersamples/example

votingapp-result

ports :
- 6060:80

deploy :

replicas : 2

docker stack deploy -c myapp.yml voting
app

→ To see the stack files created
docker stack ps

(or)

use that same docker compose file
and convert it into kcompose, it
will create necessary files, take that
files and deploy into cluster

Kubernetes :-

Memory/resource allocation :-

→ We can specify that how much amount of memory and cpus can be used by the pods.

→ In this we have two components.

(i) request

(ii) Limits

Limits :- It is the maximum amount of resources that pod can ^{ever} consume or ask for.

request :- This is the minimum amount of resources that should be given to the pod.

→ Depending on the work doing by the pod it can ask for, but if the cluster contains that hardware it can assign ~~order~~ it won't.

eg: vim requestLimits.yaml

apiVersion: v1

kind: Pod

metadata:

name: nginx-pod

Labels:

author: intelligent

Spec:

containers:

- name: mynginx

image: nginx

resources:

requests:

memory: "64Mi"

cpu: "250m"

limits:

memory: "128Mi"

cpu: "500m"

• • •

nginx-ingress

• 103.60.65.6

nginx-ingress

vim requestLimits.yaml

```
--  
apiVersion: apps/v1  
kind: Deployment  
metadata:  
  name: nginx-deployment  
  labels:  
    type: proxy  
spec:  
  replicas: 3  
  selector:  
    matchLabels:  
      type: proxy  
  template:  
    metadata:  
      name: nginx-pod  
      labels:  
        type: proxy  
    spec:  
      containers:  
        - name: mynginx  
          image: nginx  
      resources:  
        requests:
```

memory: "64Mi"

cpu: "250m"

Limits:

memory: "198Mi"

cpu: "500m"

...

Resource Requests

Node affinity :-

The pod which is created will be attracted towards that node only. i.e the pod is created on that particular node only.

→ To perform node affinity, first you have to label your slave

kubectl label nodes node-name/node-id

key = value

eg:- kubectl label nodes node-name/node-id

slave1 = intelligent

vim affint1.yaml

apiVersion: v1

kind: Pod

metadata:

name: httpd-pod

labels:

type: webserver

Spec:

containers:

- name: myhttpd

image: httpd

affinity:

nodeAffinity:

requiredDuringSchedulingIgnoredDuring

Execution:

nodeSelectorTerms:

- matchExpressions:

- key: slave

operator: In

values:

- intelligent

...

vim affinity 2.yaml

apiVersion: apps/v1

kind: Deployment

metadata:

name: httpd-deployment

Labels:

type: webserver

Spec:

replicas: 2

selector:

matchLabels:

type: webserver

template:

metadata:

name: httpd-pod

labels:

type: webserver

Spec:

containers:

- name: myhttpd

image: httpd

affinity:

nodeAffinity:

required During Scheduling Ignored During Execution

nodeSelectorTerms:

- matchExpressions :

- key : slave1

operator : In

value :

- intelligent.

Taint and Tolerations:

Taint :- It is exactly opposite of node affinity.

affinity :- If you taint a machine, the pods which

→ If you taint a machine, the pods which are created are never created on that tainted machine.

→ These pods will be created on untainted machines.

To taint
kubectl taint nodes node-name key=value

: key=value: NoSchedule

To untaint

kubectl untaint nodes node-name key=value: NoSchedule

Tolerations:-

→ To run the pod on the tainted

machines we use tolerations.

e.g:- In realtime we have many servers

and we tainted some servers for

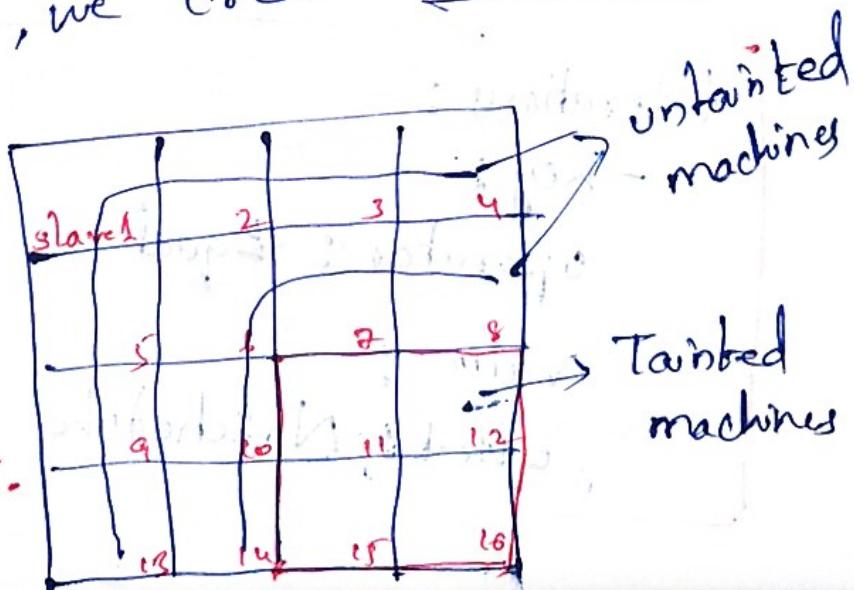
using databases. They are kept secure.

→ In this case no pods must be

created on those tainted servers.

→ Then to create pods on the tainted

servers , we create Tolerations.



→ Regular pods will run on untainted machines and database related pods will run on tainted machines by using tolerations.

→ To taint a machine

kubectl taint nodes node-name key=value:NoSchedule

rule

→ To untaint a machine

kubectl taint nodes node-name key=value:NoSchedule-

Schedule-

Syntax :- for tolerations

tolerations :

- key :

operator : Equal

value :

effect: NoSchedule

④ Create a Pod definition file to run on the tainted machine.

```
apiVersion: v1
```

```
kind: Pod
```

```
metadata:
```

```
  name: tomcat-pod
```

```
  labels:
```

```
    type: webserver
```

```
    author: intelligent
```

Spec:

```
  containers:
```

```
    - name: mytomcat
```

```
      image: tomee
```

```
  tolerations:
```

```
    - key: slave1
```

```
      operator: Equal
```

```
      value: intelligent
```

```
    effect: NoSchedule
```

* Create a tomcat-deployment file to run on tainted machines

```
apiVersion: apps/v1
```

```
kind: Deployment
```

```
metadata:
```

```
  name: tomcat-deployment
```

```
  labels:
```

```
    type: appserver
```

```
    author: intellegit
```

```
Spec:
```

```
  replicas: 3
```

```
  selector:
```

```
    matchLabels:
```

```
      type: appserver
```

```
  template:
```

```
    metadata:
```

```
      name: tomcat-pod
```

```
    labels:
```

```
      type: appserver
```

```
Spec:
```

```
  containers:
```

```
    - name: mytomcat
```

image: to me

tolerations:

- key: slave1

operator: Equal

value: intelligent

effect: NoSchedule

...

Helm:-

→ It is the package management solution of kubernetes.

→ It makes work easier
→ It creates definition files and

→ No need of creating definition files and

service files.

→ With a single command we create definition file and service file and exposed to the customer.

→ It is used for generic applications only

→ If the organisation created a customised

software then Helm won't work, then depend on definition files only.

- first install helm
- For managed k8s, helm is pre installed.
- In helm we create charts

step-2:- [mkdir helm]

[cd helm]

step-3:- Here you create helm chart

[helm create chart-name]

e.g:- helm create mynginx

step-4:-

Now go inside the chart

[cd mynginx]

There you find a folder called

values.yaml

Then edit the file, whichever you want like, image name, Service type etc.

Step - IV :-

Now to run/install the chart

`helm install release-name chart-name`

e.g:- `helm install newnginx mynginx`

→ There will be helmcharts available in artifacthub.io

Step - V :- To uninstall the chart

`helm uninstall release-name`

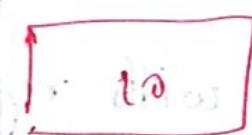
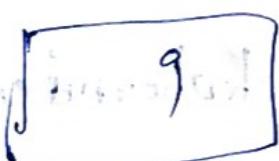
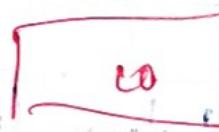
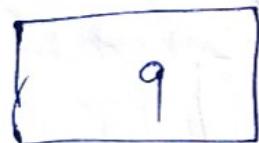
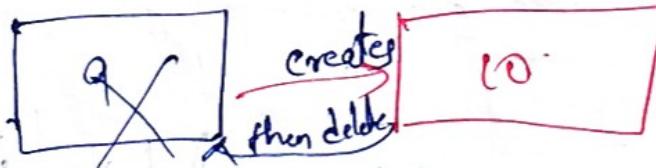
→ To see the list of helm's installed

`helm list -a`

(or)
go with regular kubernetes commands

Rolling update:-

- In docker swarm, the rolling update is done like, & let's say we have 3 replicas. It will down the first one and update it, when it comes into running condition it performs rolling update on next replica. Like wise it goes on.
- But in Kubernetes, it will create a new replica of the version you want when it comes into running condition. It deletes the older replica.



In this we have different types of strategies for rolling update, they are

i) Recreate strategy

(ii) Blue-green "

(iii) Canary "

i) Recreate strategy:-

From the command line, update nginx

kubectl set image deployment/nginx-deployment

nginx = nginx:1.25

kubectl set image deployment/deployment-name

container-name = image-name:vethin

ii) Recreate strategy:-

In this, first it deletes the older version

and creates the new replace with late version (which version you want),

In this case customer will experience down time.

vim definitionfile.yaml

apiVersion: apps/v1

kind: Deployment

metadata:

name: nginx-deployment

Labels:

type: proxy

author: intelligit

Spec:

replicas: 3

strategy:

type: Recreate

selector:

matchLabels:

type: proxy

template:

metadata:

name: nginx-pod

labels:

type: proxy

author: intelligit

Spec:

containers:

- name: nginx

image: nginx:1.24

ports:

- containerPort: 80

hostPort: 8181

Now

kubectl apply -f definitionfile.yaml

Now deployment pod is created, then
operate rolling update on it.

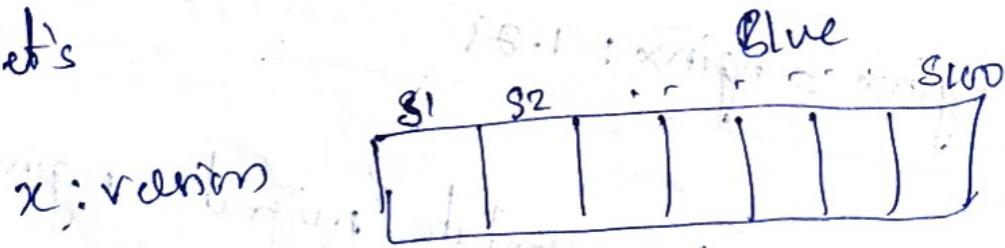
kubectl set image deployment/nginx-deployment
to nginx = nginx:1.25

This will do update/migrate nginx:1.24
to nginx:1.25

② Blue-green deployment:

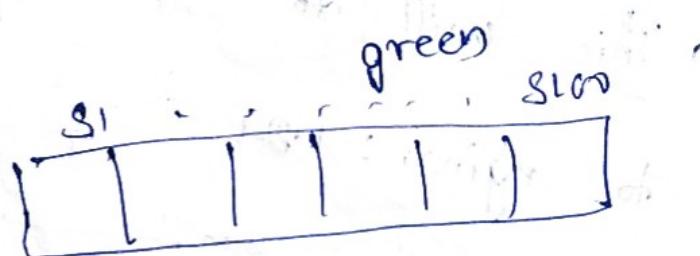
- To upgrade from one version to another version, there might be compatibility issue
- To avoid those issues we go with blue green deployment.
- Let's say nginx: 1.24 and upgrading to nginx: 1.25, it doesn't make any compatibility issues.
- If you migrating from nginx: 1.24 to nginx: 2 , bigger version, it may give compatibility errors.

→ Let's



x: version

y: version



- They are using x:version on 100 servers and it is exposed to the customers. and they are accepting it.
- Now they want to upgrade from x:version to y:version, without experiencing the downtime of customers.
- For this they won't disturb the older x:version which is running on 100 servers. They call it as "Blue".
- Now they create 100 servers and install with y:version, they call it as "green", and it also exposed to customers. When they see that y:version is behaving stable and when customers are accepting, then they delete the older x:version, and now they called y:version as blue version.

Disadvantages:

- It will increase the storage for certain time.
- Cost gets increased.

vim: blue-deployment.yaml

apiVersion: apps/v1

kind: Deployment

metadata:

name: tomcat-deployment

Labels:

type: appserver

author: intelligent

Spec:

replicas: 3

selector:

matchLabels:

type: appserver

template:

metadata:

name: tomcat-pod

Labels:

type: appserver

Spec:

containers:

- name: tomcat

image: tomcat: 9

ports:

- containerPort: 8080

hostPort: 8181

→ The above file/architecture is called as blue-deployment, because it is older version which is running.

→ Now I have to upgrade to tomcat:10
, Create same architecture definition file
and deploy. If it is working properly
then delete the older version.

vim green-deployment.yaml

apiVersion: apps/v1

kind: Deployment

metadata:

name: tomcat-deployment

labels:

type: appsource

author: intelligent

Spec:

replicas: 3

selector:

matchLabels:

type: appsource

template:

metaData:

name: tomcat-pod

labels:

type: appsource

Spec:

containers:

- name: tomcat

- image: tomcat:10

ports:

- containerPort: 8080

- hostPort: 8181

→ After deleting the older version, now
this version is called blue-deployment.
and coming versions in future are green
deployment.

Errors and debugging:

① crash loop backoff:

It indicates that the container in a pod is failing repeatedly and Kubernetes is applying a backoff strategy before attempting to restart the container again.

→ To troubleshoot and resolve this issue, you need to examine the container logs, pod status and any relevant configurations to identify and fix the underlying problem.

→ Common issues are:

(i) incorrect container configuration

(ii) missing dependencies

(iii) resource limitation

`kubectl logs pod-name`

`kubectl describe pod pod-name`

⑧ ImagePullBackOff :-
It occurs when a pod is unable to pull the specified image.
→ This error typically indicates a problem with fetching the container image specified in the pod-definition from the container registry.

Common reasons for this error :-

- ① Incorrect image name or tag;
- ② Registry authentication issues [username/password]
- ③ Image availability! - ensure that image is available or not in the registry.
- ④ Network issues!
- ⑤ Registry rate Limiting! - Some container registries have rate-limiting policies. If you exceed the limit you may encounter issues pulling images.

① Proxy configurations:

② Livenessprobe Failure:

It is a mechanism to determine whether a container within a pod is healthy and running as expected.

③ Canary deployment:

→ This is done on scaling mechanism.

let's say we have 100 pods running

in the cluster.

→ First we scale it to 95 pods, the

older version is running on 95 pods and remaining 5 pods are updated to

new version.

→ Then the user are allowed to new

version, if they think that new version is working stable, then they scale down

5 more pods and scale up to 10 pods
like wise they ~~use~~ update all the
pods running in the cluster.
old version

spec:

replicas: 95

New version

Spec:

replicas: 15

→ first create deployment definition file with
new version and scale it upto 5 replicas
and scale down the older version
deployment to 95

→ The oldie version 95 and new version
5 replicas have the desired state of

about 100% completion after 10 min

→ If it is working properly then scale down older versions to 90 and scale up new versions to 10 ... do so on upto new version is updated in all the two pods

→ Software or version format

x.y.z → patch update

major update

minor update

~~Interview~~
→ whenever we have minor and patch

update we are going with regular

rolling update process and whenever

we have major updates, we are

going with blue-green deployment in

our organisation.

Volumes :- Managing storage in the

→ In docker swarm we don't have handling volumes, i.e if a container with some data is crashed, then docker swarm will automatically creates a new container but it don't have the data. whereas in kubernetes the data will present.

→ The persistence of data will be present in kubernetes.

→ We have two types of volumes

(i) empty-dir

(ii) PersistentVolume.

(i) empty dir

Syntax:

volumeMounts :

- name : → name of the volume
- mountPath : → where it is mounted

* Create a redis pod and volume

apiVersion: v1

kind: Pod

metadata:

name: redis-pod

Labels:

type: db

Spec:

containers:

- name: myredis

image: redis:

volumeMounts:

- name: redis-volume

mountPath: /data/redis

volumes:

- name: redis-volume

emptyDir: { }

→ Now go inside the pod, in the /data/redis folder and create some files

→ Then make that pod to be crashed
i.e kill the default process of the pod
then pod will be crashed

→ To kill the pod we need ~~pod~~

PID (Process ID):

kill PID

→ To know PID → ps -ef → Linux command

→ This ps must be installed in the pod

→ To install ps, go inside the pod
and install procps software

sudo apt-get update

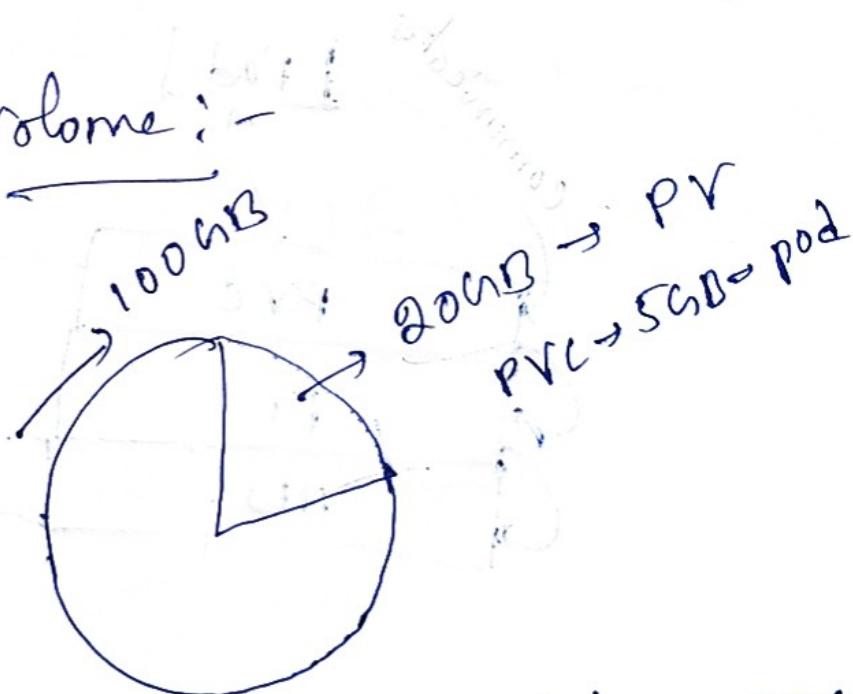
sudo apt-get update -y procps

→ After installing then pickup the PID
and kill it.

→ Automatically new pod is created, in
that there will be our previous data.

- In pods, data will be present, but when the pod crashes, and a new pod comes but the software installed in it will be lost.
- To keep software also persistent in the pod, create the docker file and keep all the necessary software you want to run in the pod, then use that image in the Pod definition file.

⑧ Persistent Volume :-

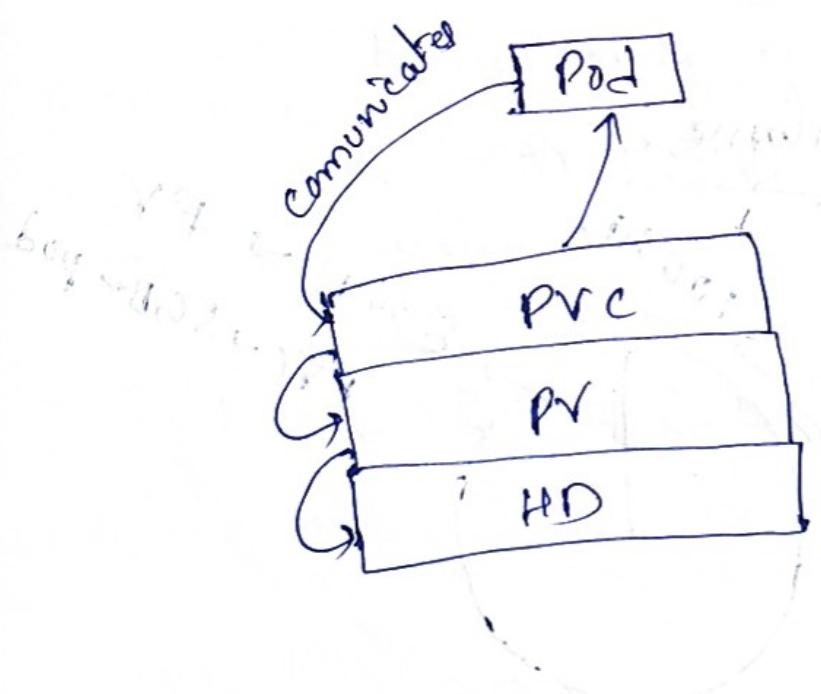


- Let's say we have a hard disk of 100 GB. From that 20 GB is reserved for volume.

it is called PersistentVolume, then when we create a pod, and if requires 5GB, then it claims the 5GB from 20GB only is called persistentvolume claim (PVC)

→ For next pod it requires 2GB, it is given from PV only.

→ Pods directly can't talk to PV, so via PVC we claim the required storage.



- ⑧ Create a PV of 20GB and PVC definitions file.

apiVersion: v1

kind: PersistentVolume

metadata:

name: task-pv-volume

Labels:

type: local

Spec:

storageClassName: manual

capacity:

storage: 20Gi

accessModes:

- ReadWriteOnce

hostPath:

path: /mnt/data

vim persistentVolumeClaim.yaml

apiVersion: v1

kind: PersistentVolumeClaim

metadata:

name: task-pr-claim

Spec:

storageClassName: manual

accessModes:

- ReadWriteOnce

resources:

requests:

storage: 5Gi

...

→ The linking of these two files done
based on

storageClassName:
accessModes:

④ Create a pod definition file to use
the prc.

vim pod-volumes.yaml

apiVersion: v1

kind: Pod
metadata:
name: nginx-pod

Labels:
author: intelligent
type: proxy

Spec:

containers:

- name: mynginx
image: nginx:latest

volumeMounts:

- name: task-pr-storage
mountPath: /usr/share/nginx/html

volumes:

- name: task-pr-storage

persistentVolumeClaim:

claimName: task-pr-claim

→ The ~~per~~ persistent volumes work only

the Pod definition files, and it won't
work on deployment, because when

pod crashed, a new pod gets created automatically with the same name, if the same name is there, with that same name it is linked with pvc.

→ But in case of deployment, pod gets created but it won't come up with same name therefore link between pvc and deployment gets unlinked and we can't recover the data.

→ To overcome this problem we go with StatefulSet definition files.

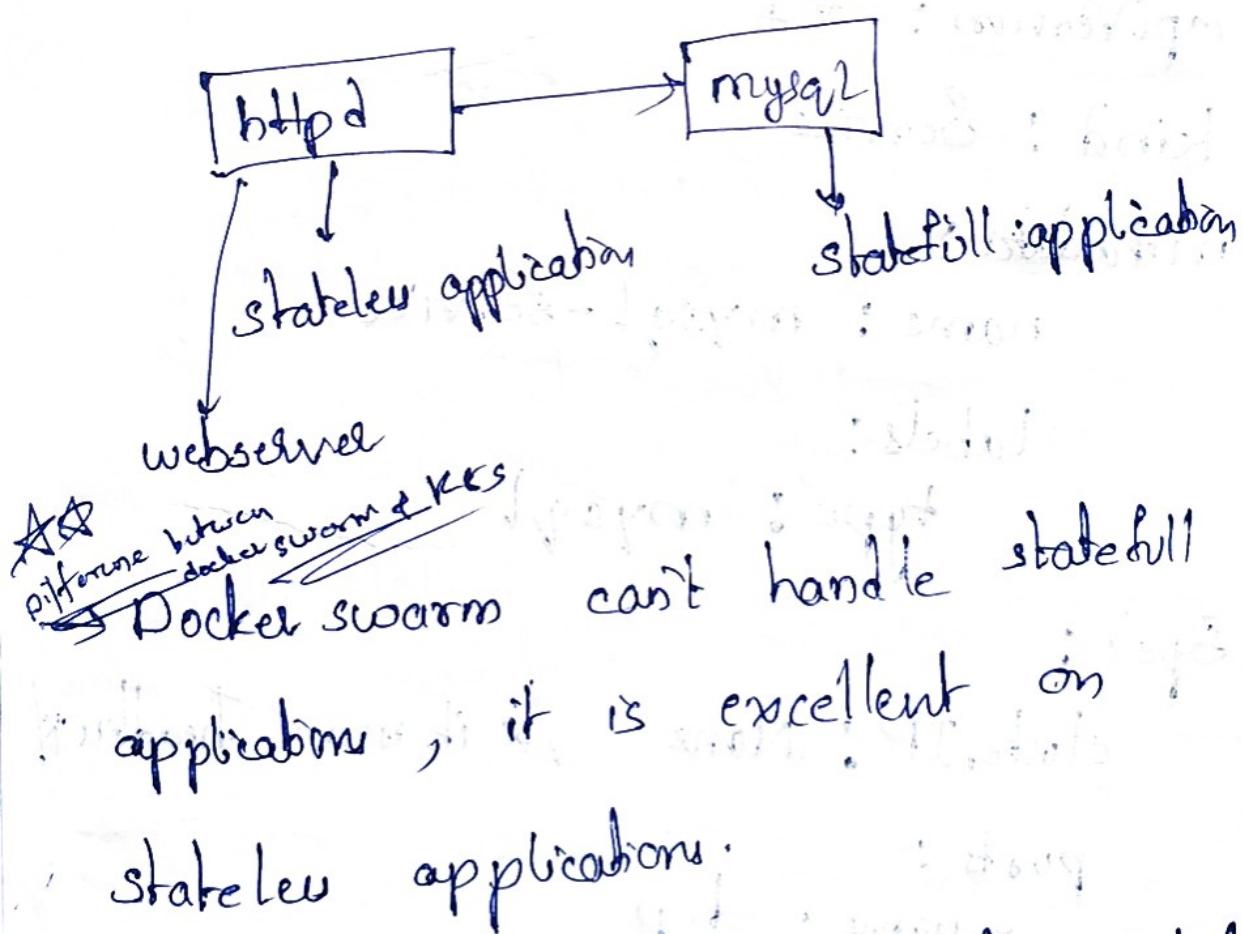
StatefulSet :-

→ If we create 3 replicas with mysql

→ It creates like mysql - 0 → master
- 1 → slave
- 2 → slave

→ The master pod continuously keeps on checking the slaves

- while deleting also it deletes in reverse way.
- Deployments are good for stateless applications whereas StatefulSets are good for statefull applications.
- Statefull applications are those applications which has persistence of data. (like databases)



- The master pods the newly created slaves whether it is linked with pre

→ The master pod always communicates with pods via a service object called or headless

* Create mysql database with 3 replicas using statefullset.

→ first create service definition file

apiVersion: v1

kind: Service

metadata:

name: mysql-service

labels:

type: mysql

Spec:

clusterIP: None /* it means headless */

ports:

- name: tcp

protocol: TCP

port: 3306

selector :
type : mysql

...

apiVersion: apps/v1

kind: StatefulSet

metadata:

name: mysql-statefull

Labels:

type: db

Spec:

replicas: 2

selector:

matchLabels:

type: mysql

serviceName: mysql-service

template:

metadata:

name: mysql-pod

Labels:

type: mysql

Spec:

containers:

- name : mydb

image: mysql

env:

- name: MYSQL_ROOT_PASSWORD

value: intelligent

volumeMounts:

- name: task-pv-storage

mountPath: /var/lib/

db: mysql

volumes:

- name: task-pv-storage

persistentVolumeClaim:

claimName: task-pv-claim

...

services -> pod -> service

Autoscaling:

→ whenever CPU, memory utilization is

happening we want to increase the

no. of pods or decrease the no. of

pods this is done based on the

object called as "Horizontal Pod AutoScaler"
→ whenever HorizontalPodAutoScaler is applied
the pods automatically increase or decrease.

yaml autoscaling.yaml

apiVersion: apps/v1

kind: Deployment

metadata:

name: php-apache

labels:

author: sudarshan

Spec:

selector:

matchLabels:

type: php-apache

template:

metadata:

name: php-apache-pod

labels:

type: php-apache

Spec:

containers:

apiVersion: v1
kind: Service
metadata:
name: php-apache

spec:
selector:
resources:
requests:

cpu: 200m

limits:

cpu: 500m

apiVersion: v1

kind: Service

metadata:

name: php-apache

labels:

type: ^{php}

author: sudarshan

Spec:

ports:

- targetPort: 80

port: 80

selector:

type: php-apache

...

The command to apply autoscaling is

kubectl autoscale deployment name-deployment
cpu --cpu-percent=50 --min=1 --max=10

→ If CPU crosses the limit of 50%

utilization it creates one pod, the pods will maintain the number between 1 to 10.

e.g. kubectl autoscale deployment php-apache
--cpu-percent=50 --min=1 --max=10

→ To check the horizontal Autoscaler

kubectl get hpa

→ To check the application for the autoscaling whether it is working or not.

→ Create a pod, and ping it continuously. Then the load increases then pods also

increases

kubectl run busybox --image
busybox --restart=Never --bin/sh -c
"while sleep 0.1; do wget -q -O -
http://php-apache; done"

→ Without giving the command always,
it shows whenever the modification is
done.

kubectl get hpa --watch

Prometheus

- we can check the status of pod, without giving commands, we have a graphical tool called as prometheus;
- Devops engineer knows the commands and he will check the status by using commands, but other people, who develop application they are interested in seeing the status of application.
- So just setup the monitoring tool called prometheus, it's a GUI.
We can have dashboards for prometheus in google, use whichever you want.
- Prometheus will capture the status from the ~~cluster~~ cluster, and the output is shown on the colour full screen this is done by Graphana.

→ Prometheus will capture the information and graphana will display it colourfull mode.

→ Now setup Prometheus with the help of helm

→ Go to <https://artifacthub.io/packages/search?query=prometheus>

→ Search for prometheus

→ And install it based on the command shown in it

→ To check whether it is added or not

`helm repo list`

→ Next to update the repository

`helm repo update`

→ After installing the prometheus, type

`kubectl get all`

it will show all the pods it created, then

in the pods search for service pod
of prometheus - grafana

service/prometheus-grafana
for this pod , it is created on clusterIP,
then change it to nodePort/LoadBalancer
to expose it to external world
→ To do this we use a command

```
kubectl patch svc prometheus-grafana  
-p '{"spec": {"type": "LoadBalancer"} }'
```

→ To check it **kubectl get all**
→ It generates public ip and access it from
the browser
→ When grafana loads, the default user
name and password's are

username : admin

password : prom-operator

→ Use any of the dashboard which are available in google, search for **grafana dashboards**.

→ Search a dashboard for prometheus in Data Source pannel.

→ Search for a dashboard called as **kubernetes cluster monitoring (via prometheus)**

kubernetes cluster (7249 - dashboard no)

→ In that you have copy id, click on it and copy that id in notepad.

→ Then go to grafana dashboard and click on '+' icon and click on 'import' dashboard

→ They will ask the id, give that number,  and click on load

helm

- To add helm repositories which are coming from third parties

helm repo add repository-name repository-url

eg:-

helm repo add bitnami <https://charts.bitnami.com/bitnami>

→ helm repo update

- To see the list of repos

helm repo list

- To remove the repository

helm repo remove repository-name

eg:- helm repo remove bitnami

→ We have downloaded bitnami repository
→ from that repository, we can use
charts whatever we required.

For example:-

* I want setup wordpress with mondb

→ first `helm search hub wordpress`

It shows the charts already present in
the repo via links/vls

To see the full vls address in the
repo.

`helm search hub wordpress --max-col-
width=0 | less`

→ first create values files, which are
required for database, like username,
password etc.

vim wordpress-values.yaml

--
wordpressUsername: admin

wordpressPassword: admin

wordpressEmail: sudarshansw7@gmail.com

wordpressFirstName: Durgesh

wordpressLastName: Sudarshan

wordpressBlogName: mywordpress.com

Service:

type: LoadBalancer

helm install wordpress bitnami/wordpress

--values=wordpress-values.yaml

→ helm install --release-name chart-name-from
repo --values=file-of-your-database-

credentials.yaml

Modifications done in helm chart

① chart.yaml

→ In this provide metadata i.e
in description :- provide some meaningful
description here.

→ It is not import.

② template:

↳ deployment.yaml

→ Edit the deployment.yaml file

→ Edit the containerPort

→ Because the port number is created

for nginx, so edit it depending
upon your application.

③ values.yaml

→ change image name, Service name,
tags.

image name :

tags : latest / give any version of your image.

Service : LoadBalance / nodePort

~~path~~

replicount:

port :

→ In managed kubernetes we don't want to create PV and Pvc, because the cloud will create them automatically.

→ We just create a deployment file and use the required components.

→ In StatefullSet just create that file and service file,

→ No need of creating PV and Pvc's

vim statefullset.yaml

apiVersion: v1

kind: Service

metadata:

name: statefullset-service

labels:

type: webservice

Spec:

clusterIP: None

ports:

- port: 80

name: web

selector:

type: NodePort

apiVersion: apps/v1

kind: StatefulSet

metadata:

name: statefullset-nginx

Labels:

type: webservice

```
spec:  
  replicas: 2  
  service_name: webservice  
  selector:  
    match_labels:  
      type: webserver  
  template:  
    metadata:  
      name: nginx-pod  
      labels:  
        type: webserver  
  spec:  
    containers:  
      - name: mynginx  
        image: nginx  
        ports:  
          - containerPort: 80  
            name: web  
    volumeMounts:  
      - name: my-volume  
        mountPath: /usr/share/nginx/html  
volumeClaimTemplates:  
  - metadata:  
    name: my-volume
```

spec:

accessModes: ["ReadWriteOnce"]

resources:

requests:

storage: 5Gi

...
allowing for multiple
and overlapping access
modes.

multiple access modes
can be combined in
multiple ways.

multiple access modes
can be combined in
multiple ways.

multiple access modes
can be combined in
multiple ways.

multiple access modes
can be combined in
multiple ways.

multiple access modes
can be combined in
multiple ways.

multiple access modes
can be combined in
multiple ways.

multiple access modes
can be combined in
multiple ways.

multiple access modes
can be combined in
multiple ways.

multiple access modes
can be combined in
multiple ways.

multiple access modes
can be combined in
multiple ways.

multiple access modes
can be combined in
multiple ways.

multiple access modes
can be combined in
multiple ways.

multiple access modes
can be combined in
multiple ways.

multiple access modes
can be combined in
multiple ways.

multiple access modes
can be combined in
multiple ways.

multiple access modes
can be combined in
multiple ways.

multiple access modes
can be combined in
multiple ways.

multiple access modes
can be combined in
multiple ways.

multiple access modes
can be combined in
multiple ways.

multiple access modes
can be combined in
multiple ways.

multiple access modes
can be combined in
multiple ways.

⑥ To install jenkins using docker container

vim dockerfile

FROM jenkins/jenkins: jdk11

USER root

RUN apt-get update

RUN apt-get install -y maven

RUN apt-get install -y git

USER jenkins

Init. containers :-

- A pod can have multiple containers running apps within it, but it can also have one or more init containers, which are run before the app containers are started.
 - Init containers always run to completion.
 - Each init container must complete successfully before the next one starts.
 - If a pod's init container fails, K8s repeatedly restarts the Pod ~~untill~~ until the init container succeeds.
 - However, if the Pod has a restartPolicy of Never, K8s does not restart the Pod.
 - Init container dies itself after execution is successful.
- * The files which are there in init containers you won't be able to access them (use for security purpose)

Sidecar containers! -

- A sidecar is a utility container in the Pod and its purpose is to support the main container.
- It must be paired with ~~any~~ one or more containers.
- It is reusable and can be paired with many type of main containers.

Usecase:

Using init and sidecar containers deploy nginx pod. secure (https) and push nginx log into aws S3.



Create a cert file

and share it to

application container

Collect logs

from app contain

er and upload

it to aws s3

bucket

YAML ngnix-dep.yaml

apiVersion: apps/v1

kind: Deployment

metadata:

name: myapp

Spec:

replicas: 2

selector:

matchLabels:

type: myapp-dep

template:

metadata:

labels: myapp-dep

Spec:

initContainers:

- name: ssl-cert-creation

image: sudarshan/openssl

command: ["bin/sh"]

args: ["-c", "mkdir -p /etc/nginx/ssl",

openssl req -newkey rsa:2048

-nodes /nginx/ssl/nginx.key -x509 -days 365

volumeMounts:

- name: ssl-cert

mountPath: /etc/nginx/ssl

Container:

- name: myapp-dep

image: nginxhttps

command: ["/home/awso-reload-nginx.sh"]

ports:

- containerPort: 4443
- containerPort: 80

livenessProbe:

httpGet:

path: /index.html

port: 80

initialDelaySeconds: 30

timeoutSeconds: 1

volumeMounts:

- mountPath: /etc/nginx/ssl

name: ssl-cert

- mountPath: /etc/nginx/conf.d

name: nginx-fifo

- mountPath: /var/log/nginx

name: logs

- name: sidecar-logs

image: sudarshan/awso

command: ["./bin/sh"]

args: ["-c", "cp /root/scripts/syncs3sh

/root/syncs3.sh; chmod 777 /root/

syncs3.sh; while true; do sh /root/

/syncs3.sh; sleep 60; done"]

volumeMounts:

- name: logs

mountPath: /var/log/nginx

- name: aws-credentials
 - mountPath: /root/.aws/config
 - subPath: config
- name: aws-credentials
 - mountPath: /root/sample/syncs3.sh
 - subPath: syncs3.sh

Volume:

- name: ssl-cert
 - emptyDir: {}
- name: log
 - emptyDir: {}
- name: aws-credentials
 - configMap:
 - name: "aws-config"
- name: nginx-files
 - configMap:
 - name: "nginx-config"

vim config-map-nginx.yaml

apiVersion: v1

kind: ConfigMap

metadata:

name: nginx-config

data:

default.conf: |

Scorer h

Paste all the certificates here.

}

config-map-aws.yaml

apiVersion: v1

kind: ConfigMap

metadata:

name: aws-config

data:

config: |

[default]

region = us-east-1

syncs3.sh: |

export AWS_CONFIG_FILE="/root/.aws/config"

export AWS_ACCESS_KEY_ID=

export AWS_SECRET_KEY=

aws s3 sync /var/log/nginx s3://bucket-name

nginx-service.yaml

apiVersion: v1

kind: Service

metadata:

name: nginxsvc

Labels:

app: nginx

Spec:

type: NodePort

ports:

- port: 80

protocol: TCP

name: http

- port: 443

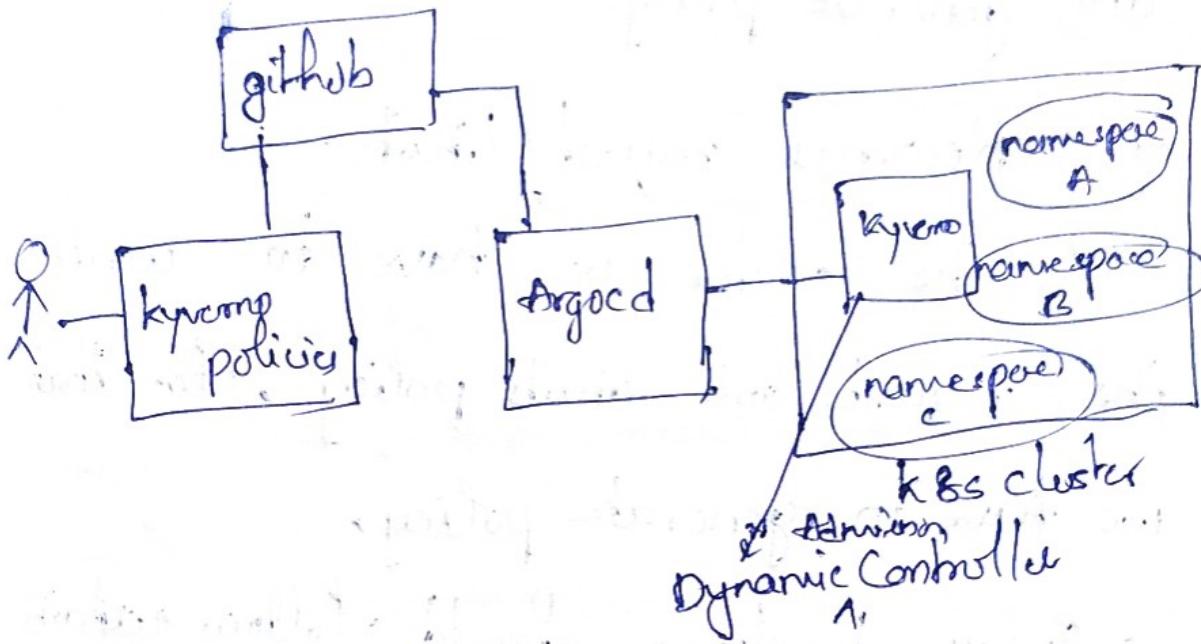
protocol: TCP

name: https

Selector:

type: myapp-dev

Enforce Automated K8s cluster security
using kyverno policy generator and argoCD



K8s governance → Managing according to
the compliance of organisation.

Compliance → set of rules

- e.g. ① In this organisation, nobody should
create a pod with latest bag.
② Nobody should create a pod the
memory size of 500MB
etc.

Admission Controller:-

It validate, mutate, verify images and generate policy.

e.g:- Resources request / limits.

In the cluster we have to write the requests and limits policy, for that we have to generate policy.

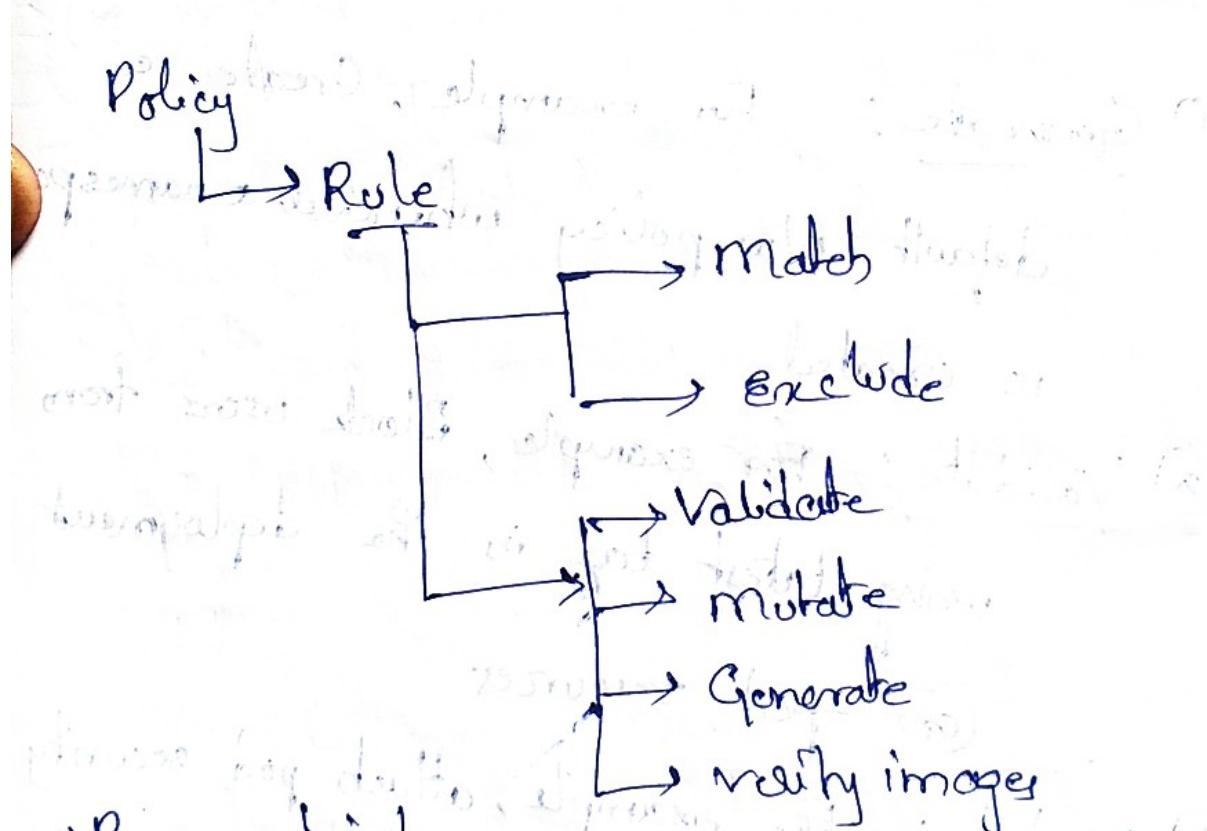
→ Every developer should follow certain rules to create a pod, then what is the guarantee that every developer is following the rule, for that we create

→ As a devops engineer we write admission controller, e.g Every pod should have resources requests / limits

- We write some yaml files with policies that kyverno will understand, then kyverno will write required webhook configuration for us.
- I worked on kyverno validate functionality to block the user, if it doesn't meet the standards of our organisation.

- ① Generate :- For example, Create a default n/w policy whenever a namespace is created
- ② Validate :- For example, Block users from using latest tag in the deployment (or) pod resources
- ③ Mutate :- For example, attach pod security policy for a pod that is created without any security policy configuration.
- ④ Verify images :- If the images used in the pod resources are properly signed, and valid images

→ A kubernetes policy is a collection of rules. Each rule consists of a match declaration, an optional exclude declaration and one of a validate, mutate, generate or verify image declaration. Each rule can contain only a single validate, mutate, generate or verify image child declaration.



→ Resource kind

- Resource name → Roles
- Labels → clusterRoles
- Annotations → Users
- Namespace → Groups
- Namespace Label → ServiceAccounts

step-①: - first install kyverno on the cluster

→ Install with helm charts/manifest files

step-②: - Install Argocd

step-③: - Copy the link of the policy which is written in yaml file and

kubectl apply -f "paste the link"

→ It will install that policy in the cluster

→ This is directly installing policy on

cluster without the argocd.

→ The admission controller which is in

the cluster will read the policy.

→ To see the kyverno pod

kubectl get pods -A | grep kyverno

→ To see the logs of kyverno pod

kubectl logs kyverno pod id

RBAC

Role Based Access Control

→ It gives authorization over cluster for users, dev team, testing team, devops team, etc. like creating/deleting pods, to see the list of pods,

→ Resources on the cluster are pods, deployments, secrets, statefulsets, Replicaset, ReplicationController etc.

→ Verbs :-

get, list, watch, create, update, patch, delete

→ Roles: We have Role, and Rolebinding, clusterRole and clusterRolebinding.

→ We have to create a role and bind it to the user

- Permissions can be applied to user, group or cluster
- Forbidden error: - On "the k8s says" that he don't have permission to do the action.

Scenario

- * Only give read access over the pods in dev namespace.

vim role.yaml

apiVersion: rbac.authorization.k8s.io/v1

kind: Role

metadata:

 namespace: dev

 name: reader

rules:

- apiGroups: [""] # indicates all api groups

 resources: ["pods"]

 verbs: ["get", "list", "watch"]

types of access

Now bind the role to a user (dev) in dev namespace

```
# vim rolebinding.yaml
```

```
apiVersion: rbac.authorization.k8s.io/v1
```

```
kind: RoleBinding
```

```
metadata:
```

```
  name: read-pods
```

```
  namespace: dev
```

```
subjects:
```

```
- kind: User
```

```
  name: dev
```

```
  apiGroup: ""
```

```
roleRef:
```

```
  kind: Role
```

```
  name: reader
```

```
  apiGroup: ""
```

```
  ...
```

→ kubectl get role -n dev

→ kubectl get rolebinding -n dev

RoleAccess

Developer → create, Read, Update

monitoring → Read

Admin → create, read, update, delete

Creating Users on k8s

Step-1: generate a private key

Step-2: create Certificate signing request

Step-3: Kubernetes certificate authority (CA)

private key

ca.key

certificate

ca.crt

These will be found in /etc/kubernetes/pki

→ Private key & certificate signing request

Certificate found in Kube config

Scenario :-

(*) Create a namespace as finance, and create a user and attach role to him.

Step-1:-

Create private to sudarshan.

[openssl genrsa -out sudarshan.key 2048]

(**) Create certificate signing request

[openssl req -new -key sudarshan.key -out sudarshan.csr]

-subj "/CN=sudarshan/O=finance"

Step-III:-
copy ca.crt, ca.key to the directory where sudarshan.key is present.

[cd /etc/kubernetes/pki]

cp ca.key .

cp ca.crt .

step-IV Now sign that key

```
openssl x509 -req -in sudarshan.csr -CA  
ca.crt -CAkey ca.key -CAcreateserial -out  
sudarshan.crt -days 365
```

step-V Now create kubeconfig file for
sudarshan

```
kubectl --kubeconfig sudarshan.kubeconfig  
get-cluster my-cluster --server https://ip:8443  
server: 8443 --certificate-authority=ca.crt
```

To view cluster service

kubectl config view

Step 6: Add user to kubeconfig file.

kubectl --kubeconfig sudarshan.kubeconfig

config set-credentials sudarshan --client-certs

set path of sudarshan.cert & client-key path

sudarshan.key

→ Now set context

kubectl --kubeconfig sudarshan.kubeconfig

config set-context sudarshan-kubernetes

--cluster my-cluster --namespace finance

for user sudarshan

→ Now go to sudarshan.kubeconfig file

Process config files

and change

current-context: sudarshan-kubernetes

→ Now ~~copy~~ sudarshan.kubeconfig to .kubeconfig

Now open copy

cp ~/.kube/config sudarshan.kubeconfig

then edit it

vim sudarshan.kubeconfig

then edit

- context
user: sudarshan

name: sudarshan-kubernetes

users:

- name: sudarshan

delete the present data in sudarshan.kubeconfig

client-certificate-data: paste sudarshan.cerbase64

client-key-data: sudarshan.pvt key

sudarshan.pkey

Convert them to base64 and paste here.